PRACTICAL COURSE WORKBOOK

AI Ethics: Principles & Practice

Evaluate technology decisions with evidence and accountability

LEVEL	Beginner
GUIDED TIME	5 hours across 12 lessons
PROJECT	a documented risk and impact assessment
SKILLS	AI Ethics, Bias, Fairness, Accountability, Transparency, Critical evaluation
EDITORIAL STATUS	Structured draft v0.9

WHAT YOU WILL BE ABLE TO DO

- Use AI Ethics appropriately in a realistic, bounded task.
- Use Bias appropriately in a realistic, bounded task.
- Use Fairness appropriately in a realistic, bounded task.
- Use Accountability appropriately in a realistic, bounded task.

WHO THIS IS FOR

Product, policy and technical learners evaluating responsible AI.

BEFORE YOU BEGIN

- No specialist background required

COURSE MAP

From foundation to finished work

AI Ethics: Principles & Practice is a structured, practical course covering AI Ethics, Bias, Fairness, Accountability, Transparency. It emphasizes explainable methods, checked work and a final outcome that can be reviewed.

MODULE	LESSONS AND FOCUS
01 AI Ethics: Ethical foundations	1. Stakeholders with AI Ethics 2. Harms with Bias 3. Rights with Fairness
02 Bias: Evidence	1. Bias with Accountability 2. Measurement with Transparency 3. Impact with AI Ethics
03 Fairness: Governance	1. Accountability with Bias 2. Documentation with Fairness 3. Oversight with Accountability
04 Accountability: Practice	1. Risk review with Transparency 2. Participation with AI Ethics 3. Monitoring with Bias

HOW TO STUDY EACH LESSON

1. Read the objective and identify the decision it supports.
2. Reproduce the worked example with the stated input.
3. Test one normal case and one boundary or failure case.
4. Complete the knowledge check without guessing.
5. Save a short evidence note: result, limitation and next action.

CAPSTONE PROJECT

Produce a reviewable AI Ethics: Principles & Practice project using AI Ethics, Bias, Fairness.

DELIVERABLE	A working AI Ethics: Principles & Practice artefact plus an evidence-based self-review.
TOOLS	A suitable AI Ethics environment, A plain-text decision log, Test data or realistic sample material

Production plan

1. Define the intended user, outcome and constraints.
2. Create the smallest complete result using AI Ethics.
3. Test one normal case, one boundary case and one failure response.
4. Revise the work from the evidence and preserve before-and-after results.
5. Prepare a concise handover containing method, limitations and next step.

Definition of done

- The output matches the stated outcome.
- Inputs and decisions are reproducible.
- Boundary and failure evidence is included.
- Limitations and responsibility considerations are explicit.

Evidence log

INPUT	Original brief, data, requirements or test material.
DECISIONS	Chosen method and one reasonable alternative.
TESTS	Normal, boundary and failure results.
LIMITATION	What the result does not establish.
REVISION	One evidence-led improvement.

REVIEW AND SOURCES

Make the work credible and repeatable

Final self-review

- Can another learner repeat the method?
- Did I test a difficult case?
- Did I avoid unsupported claims?
- Is the next action proportionate to the remaining risk?

Authoritative references

Recommendation on the Ethics of AI

UNESCO - reviewed 2026-08-21
https://www.unesco.org/en/artificial-intelligence/recommendation-ethics

AI Principles

OECD - reviewed 2026-08-21
https://www.oecd.org/en/topics/ai-principles.html

SOURCE NOTE

Product versions, laws and professional standards can change. Recheck the linked primary source before applying version-sensitive information.

NEXT IMPROVEMENT

Choose one weakness found during review and improve it before extending the AI Ethics scope.