

PRACTICAL COURSE WORKBOOK

Data Warehousing & BigQuery

Model, query and explain dependable data

LEVEL	Intermediate
GUIDED TIME	5 hours across 12 lessons
PROJECT	a working data model with validated queries
SKILLS	BigQuery, Snowflake, dbt, Data Warehouse, ETL, Critical evaluation
REVIEWED	2026-08-25

WHAT YOU WILL BE ABLE TO DO

- Use BigQuery appropriately in a realistic, bounded task.
- Use Snowflake appropriately in a realistic, bounded task.
- Use dbt appropriately in a realistic, bounded task.
- Use Data Warehouse appropriately in a realistic, bounded task.

WHO THIS IS FOR

Developers and analysts who need dependable data systems.

BEFORE YOU BEGIN

- Basic programming or spreadsheet experience

Use this guide beside the interactive lessons. Keep inputs, decisions, tests, limitations and the next improvement.

COURSE MAP

From foundation to finished work

Data Warehousing & BigQuery is a structured, practical course covering BigQuery, Snowflake, dbt, Data Warehouse, ETL. It emphasizes explainable methods, checked work and a final outcome that can be reviewed.

MODULE	LESSONS AND FOCUS
01 Data modelling	1. Entities with BigQuery 2. Relationships with Snowflake 3. Constraints with dbt
02 Querying	1. Selection with Data Warehouse 2. Joins with ETL 3. Aggregation with BigQuery
03 Reliability	1. Transactions with Snowflake 2. Indexes with dbt 3. Backups with Data Warehouse
04 Production design	1. Security with ETL 2. Performance with BigQuery 3. Operations with Snowflake

HOW TO STUDY EACH LESSON

- 1 Read the objective and identify the decision it supports.
- 2 Reproduce the worked example with the stated input.
- 3 Test one normal case and one boundary or failure case.
- 4 Complete the knowledge check without guessing.
- 5 Save a short evidence note: result, limitation and next action.

CAPSTONE PROJECT

Produce a reviewable Data Warehousing & BigQuery project using BigQuery, Snowflake, dbt.

DELIVERABLE	A working Data Warehousing & BigQuery artefact plus an evidence-based self-review.
TOOLS	A suitable BigQuery environment, A plain-text decision log, Test data or realistic sample material

Production plan

- 1 Define the intended user, outcome and constraints.
- 2 Create the smallest complete result using BigQuery.
- 3 Test one normal case, one boundary case and one failure response.
- 4 Revise the work from the evidence and preserve before-and-after results.
- 5 Prepare a concise handover containing method, limitations and next step.

Definition of done

- The output matches the stated outcome.
- Inputs and decisions are reproducible.
- Boundary and failure evidence is included.
- Limitations and responsibility considerations are explicit.

Evidence log

INPUT	Original brief, data, requirements or test material.
DECISIONS	Chosen method and one reasonable alternative.
TESTS	Normal, boundary and failure results.
LIMITATION	What the result does not establish.
REVISION	One evidence-led improvement.

REVIEW AND SOURCES

Make the work credible and repeatable

Final self-review

- Can another learner repeat the method?
- Did I test a difficult case?
- Did I avoid unsupported claims?
- Is the next action proportionate to the remaining risk?

Authoritative references

PostgreSQL Tutorial

PostgreSQL Global Development Group - reviewed 2026-08-21
<https://www.postgresql.org/docs/current/tutorial.html>

SQL Language

PostgreSQL Global Development Group - reviewed 2026-08-21
<https://www.postgresql.org/docs/current/sql.html>

SOURCE NOTE

Product versions, laws and professional standards can change. Recheck the linked primary source before applying version-sensitive information.

NEXT IMPROVEMENT

Choose one weakness found during review and improve it before extending the BigQuery scope.