

PRACTICAL COURSE WORKBOOK

Responsible AI Development

Evaluate technology decisions with evidence and accountability

LEVEL	Intermediate
GUIDED TIME	5 hours across 12 lessons
PROJECT	a documented risk and impact assessment
SKILLS	Responsible AI, Safety, Alignment, RLHF, Red-teaming, Critical evaluation
EDITORIAL STATUS	Structured draft v0.9

WHAT YOU WILL BE ABLE TO DO

- Use Responsible AI appropriately in a realistic, bounded task.
- Use Safety appropriately in a realistic, bounded task.
- Use Alignment appropriately in a realistic, bounded task.
- Use RLHF appropriately in a realistic, bounded task.

WHO THIS IS FOR

Product, policy and technical learners evaluating responsible AI.

BEFORE YOU BEGIN

- No specialist background required

This open workbook is a structured draft undergoing course-specific editorial review. It does not provide certification.

COURSE MAP

From foundation to finished work

Responsible AI Development is a structured, practical course covering Responsible AI, Safety, Alignment, RLHF, Red-teaming. It emphasizes explainable methods, checked work and a final outcome that can be reviewed.

MODULE	LESSONS AND FOCUS
01 Responsible AI: Ethical foundations	1. Stakeholders with Responsible AI 2. Harms with Safety 3. Rights with Alignment
02 Safety: Evidence	1. Bias with RLHF 2. Measurement with Red-teaming 3. Impact with Responsible AI
03 Alignment: Governance	1. Accountability with Safety 2. Documentation with Alignment 3. Oversight with RLHF
04 RLHF: Practice	1. Risk review with Red-teaming 2. Participation with Responsible AI 3. Monitoring with Safety

HOW TO STUDY EACH LESSON

- 1 Read the objective and identify the decision it supports.
- 2 Reproduce the worked example with the stated input.
- 3 Test one normal case and one boundary or failure case.
- 4 Complete the knowledge check without guessing.
- 5 Save a short evidence note: result, limitation and next action.

CAPSTONE PROJECT

Produce a reviewable Responsible AI Development project using Responsible AI, Safety, Alignment.

DELIVERABLE	A working Responsible AI Development artefact plus an evidence-based self-review.
TOOLS	A suitable Responsible AI environment, A plain-text decision log, Test data or realistic sample material

Production plan

- 1 Define the intended user, outcome and constraints.
- 2 Create the smallest complete result using Responsible AI.
- 3 Test one normal case, one boundary case and one failure response.
- 4 Revise the work from the evidence and preserve before-and-after results.
- 5 Prepare a concise handover containing method, limitations and next step.

Definition of done

- The output matches the stated outcome.
- Inputs and decisions are reproducible.
- Boundary and failure evidence is included.
- Limitations and responsibility considerations are explicit.

Evidence log

INPUT	Original brief, data, requirements or test material.
DECISIONS	Chosen method and one reasonable alternative.
TESTS	Normal, boundary and failure results.
LIMITATION	What the result does not establish.
REVISION	One evidence-led improvement.

REVIEW AND SOURCES

Make the work credible and repeatable

Final self-review

- Can another learner repeat the method?
- Did I test a difficult case?
- Did I avoid unsupported claims?
- Is the next action proportionate to the remaining risk?

Authoritative references

Recommendation on the Ethics of AI

UNESCO - reviewed 2026-08-21

<https://www.unesco.org/en/artificial-intelligence/recommendation-ethics>

AI Principles

OECD - reviewed 2026-08-21

<https://www.oecd.org/en/topics/ai-principles.html>

SOURCE NOTE

Product versions, laws and professional standards can change. Recheck the linked primary source before applying version-sensitive information.

NEXT IMPROVEMENT

Choose one weakness found during review and improve it before extending the Responsible AI scope.